

# ML Project Report

September 2022

## 1 Introduction

GDM (gestational diabetes mellitus) is a diabetes that occurs during pregnancy. About 17-63% of pregnant women with GDM develop type 2 diabetes 5-16 years after delivery; The recurrence rate of GDM in another pregnancy is as high as 52-69%. Approximately 3 to 8 percent of all pregnant women in the United States are diagnosed with gestational diabetes[1]. Diagnosing GDM (gestational diabetes mellitus) is a complex undertaking. The symptoms of diabetes are reflected in multiple indicators, which increases the difficulty of diagnosis and reduces accuracy. In this project, an efficient and interpretable algorithm was designed to diagnose diabetes based on multiple clinical data and physical examination indicators of patients.

The structure of this report is as follows: section 1 of this report is an introduction to the topic, section 2 shows the idea to formalize the above-mentioned task a machine learning problem, as well as the information on the data points, features and labels of this ML problem. In section 3, the method of the experiment is explicitly discussed. section 4 then presents the ML methods used along with their results, and section 5 is the conclusions.

## 2 Problem Formulation

Our model implements a binary classification supervised learning task to predict whether the person has diabetes through multiple features. The prediction result is whether each person has gestational diabetes, the category is expressed as an integer, and the value is 0 for false or 1 for true.

The first line of the dataset csv file is the field name, each subsequent line represents an individual, and some field names have been desensitized. The dataset contains a total of 84 feature fields, including float and int types. The first column is the individual ID number. The last column as the label column, which is the class label that needs to be predicted whether it is diseased or not. The rest columns are the clinical data and physical examination indicators of patients, which are used as features.

Except for id, SNP1, SNP2, all the features are missing in some populations. Therefore, the preprocessing is an important task before training the model.

The dataset is downloaded from Tianchi Precision Medicine Competition dataset - Artificial intelligence-assisted diabetes genetic risk prediction [2].

## 3 Methods

Here is the main methods of the experiment and the technical points.

### 3.1 Data Processing

The dataset contains 1000 samples, each sample has 84 features. Most of the features have missing values, which need to be completed with median through the method `data[col] = data[col].fillna(data[col].median())`. If a feature is missing too much, consider removing it or completing it with a special value. Here we completed the 'RBP4' and 'childbirth time' with 0.

### 3.2 Feature engineering

The project removes useless features, whose variance is less than 0.5 or low correlation with label, and only keep the first 20 features that are positively correlated and the last 15 that are negatively correlated. Taking VAR00007 as an example (Fig. 1), it can be seen that the larger the value, the higher the possibility of diabetes. The correlation is shown in the Fig. 2.




Figure 1: The relation between feature VAR00007 and the label

Here is the ablation experiment for correlation filtering. The *corr\_matrix* contain the correlation values of feature in descending order. *pos\_corr* contains top 20 features with the highest positive correlation and *neg\_corr* contains top 15 features with the most negative correlation. Other features with correlation around 0 is abandoned. The each feature number is adjusted from 0 to 20. The score is shown in the Table 1.

| number of features | 0     | 10    | 15   | 20    |
|--------------------|-------|-------|------|-------|
| test score         | 0.781 | 0.897 | 0.95 | 0.937 |

Table 1: Test score with different number of features

```

corr_matrix = data.corr()
corr_matrix = corr_matrix['label'].sort_values(ascending=False)
pos_corr = list(corr_matrix[1:20].index)
neg_corr = list(corr_matrix[-15:].index)
cols = data.columns.values
for col in cols:
    if((col not in pos_corr)
        and(col not in neg_corr)
        and (col!='label')):
        data = data.drop(col, axis=1)

```

|                      |          |                             |           |
|----------------------|----------|-----------------------------|-----------|
| label                | 1.000000 | RBP4                        | 0.066595  |
| VAR00007             | 0.384228 | SNP13                       | 0.065927  |
| SNP34                | 0.216900 | SNP46                       | 0.063351  |
| SNP                  | 0.201372 | SNP17                       | 0.056993  |
| Pregnant history     | 0.188413 | DM family history           | 0.054599  |
| age                  | 0.187000 | parity                      | 0.054041  |
| BMI before pregnancy | 0.170371 | SNP20                       | 0.053680  |
| Body condition       | 0.167354 | Delivery time               | -0.032667 |
| TG                   | 0.164522 | SNP37                       | -0.039182 |
| fat                  | 0.164144 | SNP39                       | -0.042941 |
| weight before preg   | 0.140798 | SNP48                       | -0.047821 |
| hsCRP                | 0.130591 | height                      | -0.050497 |
| wbc                  | 0.123160 | SNP41                       | -0.052168 |
| BMI class            | 0.118194 | SNP10                       | -0.058320 |
| systolic pressure    | 0.113044 | SNP28                       | -0.058641 |
| Blood presure        | 0.109015 | SNP43                       | -0.077833 |
| gravidity            | 0.086557 | SNP22                       | -0.098833 |
| ALT                  | 0.077218 | Name: label, dtype: float64 |           |

Figure 2: correlation

According to common medical knowledge, many important features are highly correlated, so we added some combined features to highlight their practical significance, such as pregnancy history, obesity, blood pressure, physical condition, and single nucleotide polymorphism SNPs. The data type of each feature is shown in the Fig. 3.

### 3.3 The construction of training and validation sets

This project uses 5-fold cross-validation to obtain average score. Because the dataset in this experiment is small for only 1000 samples, and k-fold cross validation is quite suitable for this situation. Each sample is given the opportunity to be used in the hold out set 1 time and used to train the model k-1 times, which can significantly increase the training times and make the model more likely to converge. It generally results in

| Data columns (total 37 columns): |                         |                |         |
|----------------------------------|-------------------------|----------------|---------|
| #                                | Column                  | Non-Null Count | Dtype   |
| 0                                | SNP10                   | 1000 non-null  | float64 |
| 1                                | SNP13                   | 1000 non-null  | float64 |
| 2                                | SNP17                   | 1000 non-null  | float64 |
| 3                                | SNP18                   | 1000 non-null  | float64 |
| 4                                | SNP19                   | 1000 non-null  | float64 |
| 5                                | SNP22                   | 1000 non-null  | float64 |
| 6                                | RBP4                    | 1000 non-null  | float64 |
| 7                                | age                     | 1000 non-null  | float64 |
| 8                                | gravidity               | 1000 non-null  | float64 |
| 9                                | parity                  | 1000 non-null  | float64 |
| 10                               | height                  | 1000 non-null  | float64 |
| 11                               | weight before pregnancy | 1000 non-null  | float64 |
| 12                               | BMI class               | 1000 non-null  | float64 |
| 13                               | BMI before pregnancy    | 1000 non-null  | float64 |
| 14                               | systolic pressure       | 1000 non-null  | float64 |
| 15                               | diastolic pressure      | 1000 non-null  | float64 |
| 16                               | deliver time            | 1000 non-null  | float64 |
| 17                               | VAR00007                | 1000 non-null  | float64 |
| 18                               | wbc                     | 1000 non-null  | float64 |
| 19                               | ALT                     | 1000 non-null  | float64 |
| 20                               | TG                      | 1000 non-null  | float64 |
| 21                               | hsCRP                   | 1000 non-null  | float64 |
| 22                               | SNP28                   | 1000 non-null  | float64 |
| 23                               | SNP29                   | 1000 non-null  | float64 |
| 24                               | SNP34                   | 1000 non-null  | float64 |
| 25                               | SNP37                   | 1000 non-null  | float64 |
| 26                               | DM family history       | 1000 non-null  | float64 |
| 27                               | SNP39                   | 1000 non-null  | float64 |
| 28                               | SNP41                   | 1000 non-null  | float64 |
| 29                               | SNP43                   | 1000 non-null  | float64 |
| 30                               | SNP46                   | 1000 non-null  | float64 |
| 31                               | SNP48                   | 1000 non-null  | float64 |
| 32                               | SNP52                   | 1000 non-null  | float64 |
| 33                               | SNP53                   | 1000 non-null  | float64 |
| 34                               | label                   | 1000 non-null  | int64   |
| 35                               | fat                     | 1000 non-null  | float64 |
| 36                               | blood pressure          | 1000 non-null  | float64 |
| dtypes: float64(36), int64(1)    |                         |                |         |

Figure 3: data type of features

a less biased or less optimistic estimate of the model skill than other methods, such as a simple train/test split.

Specifically, the original data is split into 5 sets, each of which accounts for 20%. Each set will be taken as the testing set for once, while the remaining sets will be used to train the model.

### 3.4 Model selection

The experiment uses random forest model. It can handle very high dimensional (many features) data and indicate which features are more important after training is complete, which is suitable for the dataset, which have totally 84 features. A large part of the features are missing in the dataset, but random forest can still maintain the accuracy. Besides, logistic regression is easy to implement and time-efficient, which is widely applied on real-world binary classification tasks; SVM and decision tree classifier are both suitable for problems with a small dataset and high-dimensional data; Naive Bayes classifier usually shows good performance when dataset is small. Hence, the experiment do comparison among these models.

Before parameter tuning, the f1\_scores of logistic regression, SVM, Naive Bayes, decision tree and random forest classifier are separately 0.667, 0.524, 0.633, 0.601, 0.674. Random forest achieves the best performance among these models. So we choose that as the final model.

### 3.5 Loss function

In classification tasks, we want to maximize both precision and recall. They reflect the predictive performance of a model in different ways:

Precision: Of all positive predictions, how many are really positive?

$$Precision = \frac{TP}{TP + FP} \quad (1)$$

Recall: Of all positive samples, how many are correctly predicted to be positive?

$$Recall = \frac{TP}{TP + FN} \quad (2)$$

Specifically, in our project, it would be most desirable that no healthy person get wrongly diagnosed while no patients are left out. In practice, however, it is not possible to maximize both of them at the same time because of the trade-off between precision and recall. Besides, there could be some problems if we choose only one of them to be the loss function: on the one hand, when precision is chosen, the model will get a perfect score by diagnosing everyone to get diabetes; on the other hand, the model tends to give negative predictions instead if recall is selected. Hence, in this project, F1 Score is chosen as the loss function, which takes both precision and recall into account.

$$F1 = \frac{2 * Precision * Recall}{Precision + Recall} \quad (3)$$

## 4 Results

Before fine tuning, the f1 score of logistic regression, SVM, Naive Bayes, decision tree, and random forest classifier are shown in Table 2.

| model    | Logistic regression | SVM   | Naive Bayes | Decision Tree | Random forest |
|----------|---------------------|-------|-------------|---------------|---------------|
| F1 score | 0.667               | 0.524 | 0.633       | 0.601         | 0.674         |

Table 2: F1 score of different model before fine tuning

As mentioned, this project uses 5-fold cross validation to train and validate model. For each round, Grid-SearchCV is adopted to find best parameters from the set of the grid of parameters:  $\{ 'n\_estimators': [10, 50, 100, 150], 'max\_depth': [3, 6, 9, 12], 'criterion': ['gini', 'entropy'] \}$ . The best parameters and F1 score of models for each round is shown in Table 3.

|            |              | model(round) |          |          |          |          |
|------------|--------------|--------------|----------|----------|----------|----------|
|            |              | model(1)     | model(2) | model(3) | model(4) | model(5) |
| parameters | criterion    | gini         | entropy  | entropy  | gini     | gini     |
|            | max_depth    | 3            | 3        | 3        | 3        | 9        |
|            | n_estimators | 150          | 100      | 50       | 10       | 50       |
| F1 score   |              | 0.705        | 0.675    | 0.715    | 0.665    | 0.645    |

Table 3: Parameters and F1 score of models for each round

Model(3), which shows the best F1 score, is selected to be the final model. Besides, average F1 score of all 5 rounds is computed to reflect the average performance of random forest classifier in this project, which is generally less biased than evaluation with simple train/test split. The average F1 score after fine tuning is 0.681.

## 5 Conclusion

We summarize here the experimental method and the experience gained from the project.

In terms of data preprocessing, some features with too many missing data should be removed or completed with special values. In feature engineering, according to the ablation experiment, removing features with zero variance and abnormality, filtering features according to feature correlation and adding combined features can significantly improve the testing accuracy. In terms of data set division, two ways can be applied to divide the data set: use function `train_test_split` and use K-fold cross-validation. When the dataset is small, using K-fold cross-validation can obtain more accurate local scores.

However, the training accuracy of the model still has the space to improve. If the dataset is large enough and the amount of features is large, a relatively good subset of features may be quickly obtained in a few iterations by a method similar to genetic algorithm (GA)[3]. This type of automatic feature screening method is less difficult to implement, and usually has a better effect than manual screening, and has always been favored by the industry. This method requires sufficient computing resources, we can try to add that method to feature engineering pro in the future on servers with more powerful hardware in the future.

## 6 Reference

- [1] John Hopkins Medicine, Gestational Diabetes Mellitus (GDM), <https://www.hopkinsmedicine.org/health/conditions-and-diseases/diabetes/gestational-diabetes>
- [2] Tianchi Precision Medicine Competition - Artificial Intelligence-Assisted Diabetes Genetic Risk Prediction, <https://tianchi.aliyun.com/competition/entrance/231638/information>
- [3] Mathworks, Genetic Algorithm Terminology, <https://se.mathworks.com/help/gads/some-genetic-algorithm-terminology.html>

# Appendix

October 9, 2022

appendix:

Code and obtained result

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: data = pd.read_csv(r'f_train.csv',encoding = 'gb2312')
```

```
[3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 85 columns):
```

| #  | Column | Non-Null Count | Dtype   |
|----|--------|----------------|---------|
| 0  | id     | 1000 non-null  | int64   |
| 1  | SNP1   | 1000 non-null  | int64   |
| 2  | SNP2   | 1000 non-null  | int64   |
| 3  | SNP3   | 948 non-null   | float64 |
| 4  | SNP4   | 989 non-null   | float64 |
| 5  | SNP5   | 935 non-null   | float64 |
| 6  | SNP6   | 987 non-null   | float64 |
| 7  | SNP7   | 997 non-null   | float64 |
| 8  | SNP8   | 995 non-null   | float64 |
| 9  | SNP9   | 995 non-null   | float64 |
| 10 | SNP10  | 995 non-null   | float64 |
| 11 | SNP11  | 983 non-null   | float64 |
| 12 | SNP12  | 985 non-null   | float64 |
| 13 | SNP13  | 978 non-null   | float64 |
| 14 | SNP14  | 953 non-null   | float64 |
| 15 | SNP15  | 970 non-null   | float64 |
| 16 | SNP16  | 977 non-null   | float64 |
| 17 | SNP17  | 974 non-null   | float64 |
| 18 | SNP18  | 970 non-null   | float64 |
| 19 | SNP19  | 959 non-null   | float64 |
| 20 | SNP20  | 910 non-null   | float64 |
| 21 | SNP21  | 461 non-null   | float64 |

|    |                                   |              |         |
|----|-----------------------------------|--------------|---------|
| 22 | SNP22                             | 460 non-null | float64 |
| 23 | SNP23                             | 462 non-null | float64 |
| 24 | RBP4                              | 89 non-null  | float64 |
| 25 | age                               | 975 non-null | float64 |
| 26 | gravidity                         | 802 non-null | float64 |
| 27 | parity                            | 802 non-null | float64 |
| 28 | height                            | 797 non-null | float64 |
| 29 | weight before pregnancy           | 797 non-null | float64 |
| 30 | BMI class                         | 796 non-null | float64 |
| 31 | BMI before pregnancy              | 796 non-null | float64 |
| 32 | systolic pressure                 | 753 non-null | float64 |
| 33 | diastolic pressure                | 754 non-null | float64 |
| 34 | deliver time                      | 185 non-null | float64 |
| 35 | Sugar screening week of pregnancy | 795 non-null | float64 |
| 36 | VAR00007                          | 990 non-null | float64 |
| 37 | wbc                               | 887 non-null | float64 |
| 38 | ALT                               | 854 non-null | float64 |
| 39 | AST                               | 746 non-null | float64 |
| 40 | Cr                                | 845 non-null | float64 |
| 41 | BUN                               | 844 non-null | float64 |
| 42 | CHO                               | 959 non-null | float64 |
| 43 | TG                                | 959 non-null | float64 |
| 44 | HDLc                              | 955 non-null | float64 |
| 45 | LDLc                              | 956 non-null | float64 |
| 46 | ApoA1                             | 934 non-null | float64 |
| 47 | ApoB                              | 934 non-null | float64 |
| 48 | Lpa                               | 934 non-null | float64 |
| 49 | hsCRP                             | 925 non-null | float64 |
| 50 | SNP24                             | 933 non-null | float64 |
| 51 | SNP25                             | 992 non-null | float64 |
| 52 | SNP26                             | 983 non-null | float64 |
| 53 | SNP27                             | 987 non-null | float64 |
| 54 | SNP28                             | 985 non-null | float64 |
| 55 | SNP29                             | 985 non-null | float64 |
| 56 | SNP30                             | 948 non-null | float64 |
| 57 | SNP31                             | 952 non-null | float64 |
| 58 | SNP32                             | 950 non-null | float64 |
| 59 | SNP33                             | 977 non-null | float64 |
| 60 | SNP34                             | 949 non-null | float64 |
| 61 | SNP35                             | 969 non-null | float64 |
| 62 | SNP36                             | 970 non-null | float64 |
| 63 | SNP37                             | 955 non-null | float64 |
| 64 | SNP38                             | 968 non-null | float64 |
| 65 | DM family history                 | 700 non-null | float64 |
| 66 | SNP39                             | 954 non-null | float64 |
| 67 | SNP40                             | 955 non-null | float64 |
| 68 | SNP41                             | 956 non-null | float64 |
| 69 | SNP42                             | 962 non-null | float64 |



|    |       |               |         |
|----|-------|---------------|---------|
| 70 | SNP43 | 969 non-null  | float64 |
| 71 | SNP44 | 957 non-null  | float64 |
| 72 | SNP45 | 968 non-null  | float64 |
| 73 | SNP46 | 929 non-null  | float64 |
| 74 | SNP47 | 953 non-null  | float64 |
| 75 | SNP48 | 978 non-null  | float64 |
| 76 | SNP49 | 983 non-null  | float64 |
| 77 | SNP50 | 981 non-null  | float64 |
| 78 | SNP51 | 981 non-null  | float64 |
| 79 | SNP52 | 987 non-null  | float64 |
| 80 | SNP53 | 956 non-null  | float64 |
| 81 | SNP54 | 483 non-null  | float64 |
| 82 | SNP55 | 483 non-null  | float64 |
| 83 | ACEID | 483 non-null  | float64 |
| 84 | label | 1000 non-null | int64   |

dtypes: float64(81), int64(4)  
memory usage: 664.2 KB

```
[4]: ### data preprocessing
# for most features: fill the missing value with median
cols = data.columns.values
for col in cols:
    if((col!='RBP4')
        and(col!='deliver time')):
        data[col] = data[col].fillna(data[col].median())
# for features missing too much: fill missing value with zero
data['RBP4'] = data['RBP4'].fillna(0)
data['deliver time'] = data['deliver time'].fillna(0)
```

```
[5]: #### feature engineering
data = data.drop(['id'], axis=1)
# remove features with zero std
remove = []
for column in data.columns:
    if data[column].std()==0:
        remove.append(column)
data = data.drop(remove, axis=1)
# remove abnormal features
unique = []
for column in data.columns:
    num = len(data[column].unique())
    if data[column].isnull().sum()!=0:
        num -= 1
    if num == 1:
        unique.append(column)
data = data.drop(unique, axis=1)
data.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000 entries, 0 to 999

Data columns (total 84 columns):

| #  | Column                            | Non-Null Count | Dtype   |
|----|-----------------------------------|----------------|---------|
| 0  | SNP1                              | 1000 non-null  | int64   |
| 1  | SNP2                              | 1000 non-null  | int64   |
| 2  | SNP3                              | 1000 non-null  | float64 |
| 3  | SNP4                              | 1000 non-null  | float64 |
| 4  | SNP5                              | 1000 non-null  | float64 |
| 5  | SNP6                              | 1000 non-null  | float64 |
| 6  | SNP7                              | 1000 non-null  | float64 |
| 7  | SNP8                              | 1000 non-null  | float64 |
| 8  | SNP9                              | 1000 non-null  | float64 |
| 9  | SNP10                             | 1000 non-null  | float64 |
| 10 | SNP11                             | 1000 non-null  | float64 |
| 11 | SNP12                             | 1000 non-null  | float64 |
| 12 | SNP13                             | 1000 non-null  | float64 |
| 13 | SNP14                             | 1000 non-null  | float64 |
| 14 | SNP15                             | 1000 non-null  | float64 |
| 15 | SNP16                             | 1000 non-null  | float64 |
| 16 | SNP17                             | 1000 non-null  | float64 |
| 17 | SNP18                             | 1000 non-null  | float64 |
| 18 | SNP19                             | 1000 non-null  | float64 |
| 19 | SNP20                             | 1000 non-null  | float64 |
| 20 | SNP21                             | 1000 non-null  | float64 |
| 21 | SNP22                             | 1000 non-null  | float64 |
| 22 | SNP23                             | 1000 non-null  | float64 |
| 23 | RBP4                              | 1000 non-null  | float64 |
| 24 | age                               | 1000 non-null  | float64 |
| 25 | gravidity                         | 1000 non-null  | float64 |
| 26 | parity                            | 1000 non-null  | float64 |
| 27 | height                            | 1000 non-null  | float64 |
| 28 | weight before pregnancy           | 1000 non-null  | float64 |
| 29 | BMI class                         | 1000 non-null  | float64 |
| 30 | BMI before pregnancy              | 1000 non-null  | float64 |
| 31 | systolic pressure                 | 1000 non-null  | float64 |
| 32 | diastolic pressure                | 1000 non-null  | float64 |
| 33 | deliver time                      | 1000 non-null  | float64 |
| 34 | Sugar screening week of pregnancy | 1000 non-null  | float64 |
| 35 | VAR00007                          | 1000 non-null  | float64 |
| 36 | wbc                               | 1000 non-null  | float64 |
| 37 | ALT                               | 1000 non-null  | float64 |
| 38 | AST                               | 1000 non-null  | float64 |
| 39 | Cr                                | 1000 non-null  | float64 |
| 40 | BUN                               | 1000 non-null  | float64 |
| 41 | CHO                               | 1000 non-null  | float64 |
| 42 | TG                                | 1000 non-null  | float64 |

|    |                   |               |         |
|----|-------------------|---------------|---------|
| 43 | HDLc              | 1000 non-null | float64 |
| 44 | LDLc              | 1000 non-null | float64 |
| 45 | ApoA1             | 1000 non-null | float64 |
| 46 | ApoB              | 1000 non-null | float64 |
| 47 | Lpa               | 1000 non-null | float64 |
| 48 | hsCRP             | 1000 non-null | float64 |
| 49 | SNP24             | 1000 non-null | float64 |
| 50 | SNP25             | 1000 non-null | float64 |
| 51 | SNP26             | 1000 non-null | float64 |
| 52 | SNP27             | 1000 non-null | float64 |
| 53 | SNP28             | 1000 non-null | float64 |
| 54 | SNP29             | 1000 non-null | float64 |
| 55 | SNP30             | 1000 non-null | float64 |
| 56 | SNP31             | 1000 non-null | float64 |
| 57 | SNP32             | 1000 non-null | float64 |
| 58 | SNP33             | 1000 non-null | float64 |
| 59 | SNP34             | 1000 non-null | float64 |
| 60 | SNP35             | 1000 non-null | float64 |
| 61 | SNP36             | 1000 non-null | float64 |
| 62 | SNP37             | 1000 non-null | float64 |
| 63 | SNP38             | 1000 non-null | float64 |
| 64 | DM family history | 1000 non-null | float64 |
| 65 | SNP39             | 1000 non-null | float64 |
| 66 | SNP40             | 1000 non-null | float64 |
| 67 | SNP41             | 1000 non-null | float64 |
| 68 | SNP42             | 1000 non-null | float64 |
| 69 | SNP43             | 1000 non-null | float64 |
| 70 | SNP44             | 1000 non-null | float64 |
| 71 | SNP45             | 1000 non-null | float64 |
| 72 | SNP46             | 1000 non-null | float64 |
| 73 | SNP47             | 1000 non-null | float64 |
| 74 | SNP48             | 1000 non-null | float64 |
| 75 | SNP49             | 1000 non-null | float64 |
| 76 | SNP50             | 1000 non-null | float64 |
| 77 | SNP51             | 1000 non-null | float64 |
| 78 | SNP52             | 1000 non-null | float64 |
| 79 | SNP53             | 1000 non-null | float64 |
| 80 | SNP54             | 1000 non-null | float64 |
| 81 | SNP55             | 1000 non-null | float64 |
| 82 | ACEID             | 1000 non-null | float64 |
| 83 | label             | 1000 non-null | int64   |

dtypes: float64(81), int64(3)  
memory usage: 656.4 KB


```
[6]: # filter high correlation
pd.set_option('display.max_rows', None)
corr_matrix = data.corr()
```

```
corr_matrix = corr_matrix['label'].sort_values(ascending=False)
```

```
[7]: pos_corr = list(corr_matrix[1:20].index)
neg_corr = list(corr_matrix[-15:].index)
cols = data.columns.values
for col in cols:
    if((col not in pos_corr) and (col not in neg_corr) and (col != 'label')):
        data = data.drop(col, axis=1)
```

```
[8]: # examples about the relation between features and label
pd.crosstab(data.VAR00007,data.label).plot.bar(stacked = True)
```

```
[8]: <AxesSubplot:xlabel='VAR00007'>
```



```
[9]: data['SNP34'].value_counts()
```

```
[9]: 1.0    541
      2.0    319
      3.0    140
      Name: SNP34, dtype: int64
```

```
[10]: data['SNP37'].value_counts()
```

```
[10]: 1.0    802
      2.0    150
      3.0     48
      Name: SNP37, dtype: int64
```

```
[11]: data['fat'] = data['BMI before pregnancy']+data['weight before pregnancy']
      data['blood pressure'] = data['systolic pressure']+data['diastolic pressure']
      data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 37 columns):
```

| #  | Column                  | Non-Null Count | Dtype   |
|----|-------------------------|----------------|---------|
| 0  | SNP10                   | 1000 non-null  | float64 |
| 1  | SNP13                   | 1000 non-null  | float64 |
| 2  | SNP17                   | 1000 non-null  | float64 |
| 3  | SNP18                   | 1000 non-null  | float64 |
| 4  | SNP19                   | 1000 non-null  | float64 |
| 5  | SNP22                   | 1000 non-null  | float64 |
| 6  | RBP4                    | 1000 non-null  | float64 |
| 7  | age                     | 1000 non-null  | float64 |
| 8  | gravidity               | 1000 non-null  | float64 |
| 9  | parity                  | 1000 non-null  | float64 |
| 10 | height                  | 1000 non-null  | float64 |
| 11 | weight before pregnancy | 1000 non-null  | float64 |
| 12 | BMI class               | 1000 non-null  | float64 |
| 13 | BMI before pregnancy    | 1000 non-null  | float64 |
| 14 | systolic pressure       | 1000 non-null  | float64 |
| 15 | diastolic pressure      | 1000 non-null  | float64 |
| 16 | deliver time            | 1000 non-null  | float64 |
| 17 | VAR00007                | 1000 non-null  | float64 |
| 18 | wbc                     | 1000 non-null  | float64 |
| 19 | ALT                     | 1000 non-null  | float64 |
| 20 | TG                      | 1000 non-null  | float64 |
| 21 | hsCRP                   | 1000 non-null  | float64 |
| 22 | SNP28                   | 1000 non-null  | float64 |
| 23 | SNP29                   | 1000 non-null  | float64 |
| 24 | SNP34                   | 1000 non-null  | float64 |

|    |                   |               |         |
|----|-------------------|---------------|---------|
| 25 | SNP37             | 1000 non-null | float64 |
| 26 | DM family history | 1000 non-null | float64 |
| 27 | SNP39             | 1000 non-null | float64 |
| 28 | SNP41             | 1000 non-null | float64 |
| 29 | SNP43             | 1000 non-null | float64 |
| 30 | SNP46             | 1000 non-null | float64 |
| 31 | SNP48             | 1000 non-null | float64 |
| 32 | SNP52             | 1000 non-null | float64 |
| 33 | SNP53             | 1000 non-null | float64 |
| 34 | label             | 1000 non-null | int64   |
| 35 | fat               | 1000 non-null | float64 |
| 36 | blood pressure    | 1000 non-null | float64 |

dtypes: float64(36), int64(1)  
memory usage: 289.2 KB

```
[13]: from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import KFold
from sklearn.metrics import f1_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
# as the dataset in this experiment is small, we adapt 5-fold cross validation
# selected model: random forest
# model for comparison: logistic regression, svm, decision tree, naive bayes

# for random forest
gsrf_score_sum = 0
gsrf_best_score = 0
global gsrf, best_gsrf
# for logistic regression
lr_score_sum = 0
lr_best_score = 0
global gslr, best_gslr
# for svm
svm_score_sum = 0
svm_best_score = 0
global svm, best_svm
# for decision tree
dt_score_sum = 0
dt_best_score = 0
global dt, best_dt
# for naive_bayes
gnb_score_sum = 0
gnb_best_score = 0
```

```

global gnb, best_gnb
# compare the performace of random forest classifer, logistic regression, sum,
↳ decision tree, and naive_bayes without fine-tune
kf =KFold(n_splits=5)
for train_index, test_index in kf.split(data):
    train, test = data.iloc[train_index], data.iloc[test_index]
    x_train = train.drop(['label'], axis=1)
    y_train = train['label']
    x_test = test.drop(['label'], axis=1)
    y_test = test['label']
    # for random forest
    gsrf = RandomForestClassifier(random_state=1)
    gsrf.fit(x_train,y_train)
    y_predict = gsrf.predict(x_test)
    score = f1_score(y_test, y_predict, average='micro')
    if score > gsrf_best_score:
        gsrf_best_score = score
        best_gsrf = gsrf
    gsrf_score_sum = gsrf_score_sum + score
    # for logistic regression
    lr = LogisticRegression(penalty='l2', max_iter=10000)
    lr.fit(x_train, y_train)
    y_predict = lr.predict(x_test)
    score = f1_score(y_test, y_predict, average='micro')
    if score > lr_best_score:
        lr_best_score = score
        best_lr = lr
    lr_score_sum = lr_score_sum + score
    # for sum
    svm = SVC()
    svm.fit(x_train, y_train)
    y_predict = svm.predict(x_test)
    score = f1_score(y_test, y_predict, average='micro')
    if score > svm_best_score:
        svm_best_score = score
        best_svm = svm
    svm_score_sum = svm_score_sum + score
    # for decision tree
    dt = DecisionTreeClassifier(random_state=1)
    dt.fit(x_train, y_train)
    y_predict = dt.predict(x_test)
    score = f1_score(y_test, y_predict, average='micro')
    if score > dt_best_score:
        best_dt_score = score
        best_dt = dt
    dt_score_sum = dt_score_sum + score
    # for naive_bayes

```

```

gnb = GaussianNB()
gnb.fit(x_train, y_train)
y_predict = gnb.predict(x_test)
score = f1_score(y_test, y_predict, average='micro')
if score > gnb_best_score:
    best_gnb_score = score
    best_gnb = gnb
gnb_score_sum = gnb_score_sum + score
gsrf_score_sum/=5
lr_score_sum/=5
svm_score_sum/=5
dt_score_sum/=5
gnb_score_sum/=5
print('logistic regression f1 score: ', lr_score_sum)
print('svm f1 score:', svm_score_sum)
print('naive bayes classifier f1 score:', gnb_score_sum)
print('decision tree classifier f1 score:', dt_score_sum)
print('random forest classifier f1 score: ', gsrf_score_sum)

```

```

logistic regression f1 score:  0.667
svm f1 score: 0.524
naive bayes classifier f1 score: 0.633
decision tree classifier f1 score: 0.601
random forest classifier f1 score:  0.674

```

```

[14]: # random forest classifier shows the best performance
      # do fine tuning with grid search
      gsrf_score_sum = 0
      gsrf_best_score = 0
      kf = KFold(n_splits=5)
      for train_index, test_index in kf.split(data):
          train, test = data.iloc[train_index], data.iloc[test_index]
          x_train = train.drop(['label'], axis=1)
          y_train = train['label']
          x_test = test.drop(['label'], axis=1)
          y_test = test['label']
          # for random forest
          rf = RandomForestClassifier(random_state=1)
          prf = [{'n_estimators':[10,50,100,150], 'max_depth':[3,6,9,12], 'criterion':
↳ ['gini', 'entropy']}]
          gsrf = GridSearchCV(estimator = rf, param_grid = prf, scoring =_
↳ 'accuracy', cv = 2)
          gsrf.fit(x_train, y_train)
          print(gsrf.best_params_)
          y_predict = gsrf.predict(x_test)
          score = f1_score(y_test, y_predict, average='micro')
          print(score)

```



```
if score > gsrf_best_score:
    gsrf_best_score = score
    best_gsrf = gsrf
    gsrf_score_sum = gsrf_score_sum + score
gsrf_score_sum/=5
print('random forest classfier f1 score after fine tuning: ', gsrf_score_sum)
```

```
{'criterion': 'gini', 'max_depth': 3, 'n_estimators': 150}
0.705
{'criterion': 'entropy', 'max_depth': 3, 'n_estimators': 100}
0.675
{'criterion': 'entropy', 'max_depth': 3, 'n_estimators': 50}
0.715
{'criterion': 'gini', 'max_depth': 3, 'n_estimators': 10}
0.665
{'criterion': 'gini', 'max_depth': 9, 'n_estimators': 50}
0.645
random forest classfier f1 score after fine tuning: 0.6809999999999999
```

[ ]: